hw4 5525

Kexin Feng 5165462 fengx463@umn.edu

December 2018

1 Prob 2

The tensorboard output has been stored in directory tb_prob2. You can use the following command to get the url and open it with browser.

 $tensorboard - logdir=tb \cdot prob2$

The test accuracy is 0.8983

The training performance:

Average loss epoch 0: 1.2891434678386697 Average loss epoch 1: 0.7327894783798076 Average loss epoch 2: 0.6007488074697259 Average loss epoch 3: 0.5360953472988866 Average loss epoch 4: 0.49808518446130906 Average loss epoch 5: 0.470791361340276 Average loss epoch 6: 0.4516749320191381 Average loss epoch 7: 0.4355761824926852

Average loss epoch 8: 0.42291591550920393 Average loss epoch 9: 0.4135298542765193 Total time: 23.235044956207275 seconds

Optimization Finished!

Accuracy 0.8983

2 Prob 3

For this question, instead of GradientDescent I used AdamOptimizer, because GradientDescent doesn't converge fast enough. But if you want to use GradientDescent, you can use the follow command line:

\$python ConvNetTemplate.py default

The convolutional neural network is based on the confolutional product of the 'image' data an the filter.

$$S[i,j] = I * K = \sum_{m} \sum_{n} I[m,n]K[i-m,j-n]$$
 (1)

A CNN layer is very similar to a layer of ordinary Neural Network in a sense that they both, in each channel, capture certain features from the input layer, and that they both can increase or decrease the dimension. The difference is that ordinary neural network doesn't share weight parameters and the parameter space is big; it's less explainable other than seeing it as a regression fitting. The convolutional network, however, share the filters, which has much less degrees of freedom. And when a region of image has similar pattern as filter, the corresponding

The tensorboard output has been stored in directory tb_prob3. You can use the following command to get the url and open it with browser.

\$tensorboard --logdir=tb_prob3

The test accuracy is 0.918

The training performance: Average loss at step 240: 2453.8

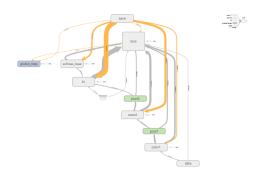


Figure 1: CNN for prob 3

Average loss at step 250: 1825.0 Average loss at step 260: 1893.2 Average loss at step 270: 1855.9 Average loss at step 280: 1733.9 Average loss at step 290: 1502.8 Average loss at step 300: 1827.1 Average loss at step 310: 1652.2 Average loss at step 320: 1835.9 Average loss at step 330: 1649.5 Average loss at step 340: 1539.7 Average loss at step 350: 1399.6 Average loss at step 360: 1543.7 Average loss at step 370: 1650.7 Average loss at step 380: 1370.9 Average loss at step 390: 1449.2 Average loss at step 400: 1470.6 Average loss at step 410: 1079.3 Average loss at step 420: 1088.2 Optimization Finished!

Total time: 282.3850169181824 seconds

Accuracy 0.918

3 Prob 4

This problem, I did task one. To increase the accuracy, I chose RNN, more specifically, bidirectional LSTM and gru. The reason of choosing this model is because in a digit graph, each digit are not totally independent from each other; they have statistical correlation for a certain hand written number. So using LSTM or gru, we can treat rows as time sequences, so this time correlation information can be captured and the accuracy can be bigger. The bidirectional lstm was chosen because a simple lstm has the problem of forgetting the beginning part. So bidirectional lstm can capture the information of both the beginning and the ending part.

LSTM is gated RNN. One hidden unit in an LSTM cell is shown in the figure 4. It has a forget gate, an input gate and output gate, all of which output a value between 0 and 1, controlling how much information can pass through. It also has two channels, the output channel and memory channel. The memory channel is what used to store the long term memory and constitute part of the output. The output channel takes input at each step and outputs the result to the user and recurrently to the next-step input.

This is a classification problem. Each data is assigned to one of ten distinguished classes, corresponding to ten numbers. The loss function is the cross entropy.

$$L(y(x_i), t_i) = \sum_{n, i} t_{n, i} \log y_{n, i}$$
 (2)

The optimizer is Adam Optimizer.

I tried two models, bidirectional LSTM and GRU. Their results behave similar to each other and both reach 0.97! It's shown in Fig. 2 and Fig.3

Bidirectional LSTM test accuracy: 0.9765625

GRU test accuracy: 0.9765625

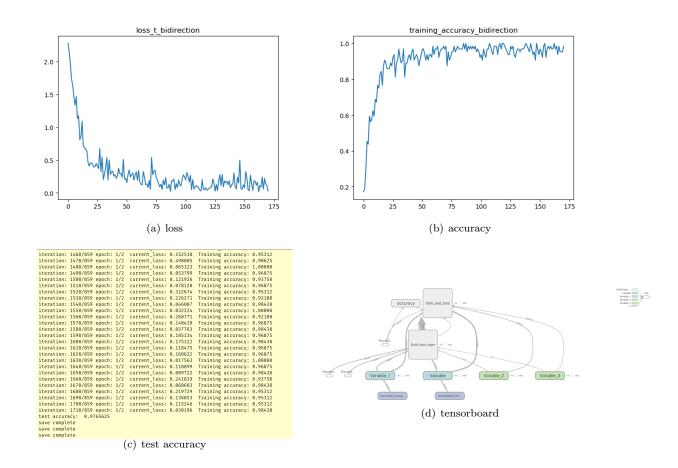


Figure 2: Bidirectional LSTM.

4 Prob 5

In this problem, I treated the digit rows as a time series of vectors of dimension 28. So the number of hidden units of LSTM cell is set to be 28. The model used in this problem, Long Short Term Memory(LSTM), is a kind of recurrent neural network(RNN), which takes sequence as its input, one element at a time.

The detail mechanism and the equations have been described above in problem 4.

Then I treated this as a classification problem. Each data is assigned to one of ten distinguished classes, corresponding to ten numbers. The loss function is the cross entropy.

$$L(y(x_i), t_i) = \sum_{n, i} t_{n, i} \log y_{n, i}$$
(3)

The optimizer is Adam Optimizer. Test accuracy: 0.9921875

The training loss decrease, accuracy increase and the test accuracy is show below in figure 5

References

[1] http://colah.github.io/posts/2015-08-Understanding-LSTMs/

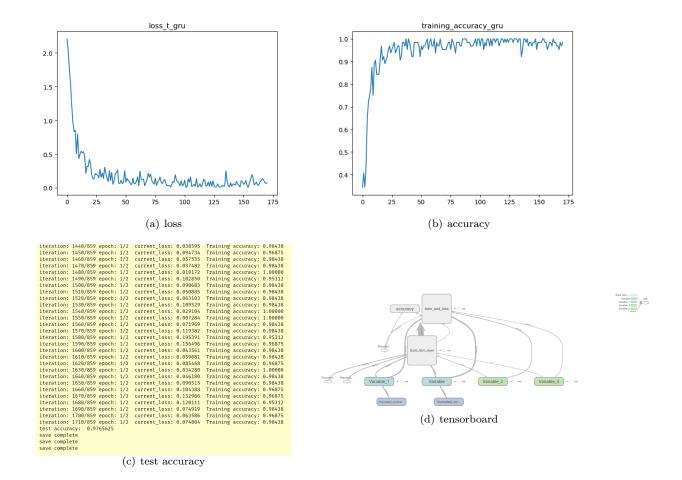


Figure 3: GRU

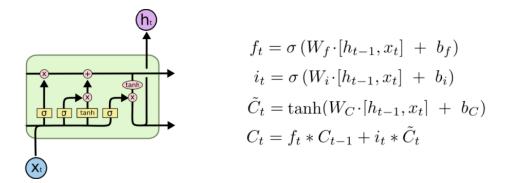


Figure 4: LSTM cell and equation. [1]

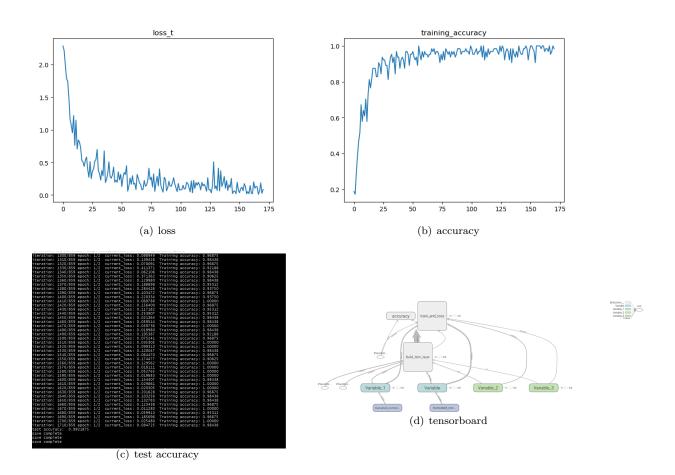


Figure 5: LSTM result.